


# Bot-tleship

Final Presentation | Artificial Intelligence | Monsoon '24



# Refresher

- Agent that can play against a human so people can play battleship against a bot
- We chose some baselines
  - Random
  - Human-like
- Previous work includes algorithms like:
  - Monte Carlo Tree Search
  - Probability-based
  - RL algos
- We also wanted to compare the bots against each other to find the best performing one
- Expectations:
  - The system can complete the game in as few moves as possible
  - Reasoning based on current state and hits on board - where to target next shot for highest likelihood of hit
  - Several heuristics exist to optimize for battleship - does the AI learn these?

# Baselines

## Random Bot

- Chooses un-attacked squares randomly
- On a 10x10 board with 5 ships, takes average of 96 moves to complete the game (destroy all ships), SD ~4
- Meaning the random bot takes almost the maximum number of moves to win



# Baselines

## Heuristics-based human-like bot

- Follows the algorithm used most commonly by human players:
  1. **Hunt Mode:** Randomly attacks squares on the board until it hits a ship.
  2. **Target Mode:** After a hit, attacks adjacent squares to determine the ship's orientation (horizontal or vertical).
  3. **Destroy Mode:** Focuses attacks along the determined orientation to sink the ship completely.
  4. **Reset:** Once a ship is sunk, returns to Hunt Mode to find the next ship.
- On a 10x10 board with 5 ships, takes an average of 76.23 moves to complete the game



# AI-Based Bots

We implemented three bots:

- MCTS-Bot
- DQN-Bot
- PPO-Bot



# Monte-Carlo Tree Search

- We have a very large tree space, so we use a probabilistic approach
- The algorithm has four steps:
  - Selection: We select a child to explore
  - Expansion: We expand the selected node by choosing a new child
  - Simulation: We simulate a potential win and assign a reward based on the validity
  - Back propagation: We send the reward back up the tree




Figure from Chaslot (2006)

# Choosing the “Best Child”

- After all the simulations are done we choose the best child
- The “best” child is given by the UCT formula:  $UCT(node_i) = \frac{w_i}{n_i} + c\sqrt{\frac{\ln N_i}{n_i}}$
- This formula balances the “wins” that a child has and ensuring that all the expanded children are given a chance to be explored



# A Problem

- Battleship is a POMDP, we don't know what the outcome will be of a move, unlike say Chess
- Solution:
  - We use heuristics like a heatmap to help the bot choose moves (i.e. give more weight to explore moves near known hits)
  - Instead of deciding whether there was a win, we simulate that all the children were hits, then evaluate if the final position is "valid" and based on this give a reward
  - Reward function:  $(\text{valid squares} / \text{ship squares}) * (\text{valid ships} / \text{total ships})^2$
- We backpropagate this reward after each simulation





# Q-Learning

- Q learning uses a Q-table which contains all possible states, and the reward when going from State-A to State-B
- It isn't possible to have a large table for a game like battleship, where there are so many possible states and a probabilistic result
- So we use Deep Q-Learning which does not require the table



# Deep Q-Network

- A DQN based AI attempts to “approximate” the Q-table by observing many “episodes” of events and regressing a manifold based on them
  - Reinforcement learning is used to train the neural network used for approximation
- Useful for sequential decision-making problems
- Balances exploration (trying random shots) and exploitation (targeting a ship)
- We trained a CNN to calculate the expected q-values given the state and an action



# Why CNNs?

- CNNs can extract spatial patterns (i.e. grid locations of ships)
- The convolutional filters focus on small patches of the board (5x5)
  - Can capture local patterns like clusters of hits or ship alignments
  - Fully connected layers flatten the data, losing spatial information



# DQN Architecture

- State-Action Value Function:  $Q(\text{state}, \text{action})$ 
  - Given the current state, what is the expected award for *action*?
  - Approximated by CNN
- Replay buffer
  - Stores past experiences - (*state, action, reward, next state, done*)
  - During training, the agent samples a random batch of past experiences to train the CNN
  - Random sampling ensures diverse training data
- Target Net
  - The q-values and *next state* calculated by the main network are fed to a target net (updated less frequently)
  - Used to calculate the expected future reward and train the main network

# CNN Architecture

- Our CNN architecture has 7 layers:
  1. Conv1: Extracts spatial patterns over the whole board (kernel size = 10).
  2. Max Pooling: Reduces spatial dimensions, highlights key features.
  3. Conv2: Captures finer spatial details (kernel size = 5).
  4. 3 Dense Layers: Encodes the spatial features into q-values.
    - a. Outputs a q-value for each grid position (100 values)




# Training the DQN

- Attempted to use Gymnasium and stable\_baselines with a custom environment for battleship
  - Did not support masking invalid actions, so tried setting reward for invalid shots very low
  - Despite this, bot was not learning to choose valid moves
  - Had to create own environment and training setup
- Could only train over 3000 games because of low computation speed
  - Took 6 hours
- 3000 steps was not enough for the bot to reach the performance we expected
  - Still showed some improvement over time, and better performance than RandomBot
  - With more training, we would expect performance closer to MCTS and PPO



# Training the DQN



- Slope for rewards: 0.00089
- P-value for the slope of rewards: 0.003 (statistically significant)

- Slope for shots: -0.00078
- P-value for the slope of shots: 0.009

# Proximal Policy Optimization (PPO)

- Learns policy directly in an on-policy manner
- It tries to keep the new policies close enough to the old policies so as to not overshoot
- Clipped surrogate objective comes into play that stops new policy from going too far away, controlled by a hyperparameter  $\epsilon$  (this is mainly what handles exploration-vs-exploitation)






# PPO Architecture

- Actor-Critic Framework:
  - Actor: Proposes actions based on the current policy
  - Actor Network: Probability distribution over actions
  - Critic: Evaluates the actions by estimating the value function
  - Critic Network: Value estimate for given state
- Reward and Penalty formulation:
  - If attacks previously attacked square : **Large penalty**
  - If gets a hit : **Equivalently large reward**
  - If misses : **Small penalty**




# PPO Training


- PPO interacts with the environment (a discrete multi-dimensional space)
  - Actor updates its policy to choose actions based on feedback of the rewards from the critic
  - Policy is clipped to avoid instability and overshooting
  - Advantages measure how good a policy is compared to the baseline
  - Entropy term encourages exploration
  - Discount factor ensures future rewards are discounted
- 

# PPO Training

Train/sum\_rew\_avg



Train/traj\_len\_avg

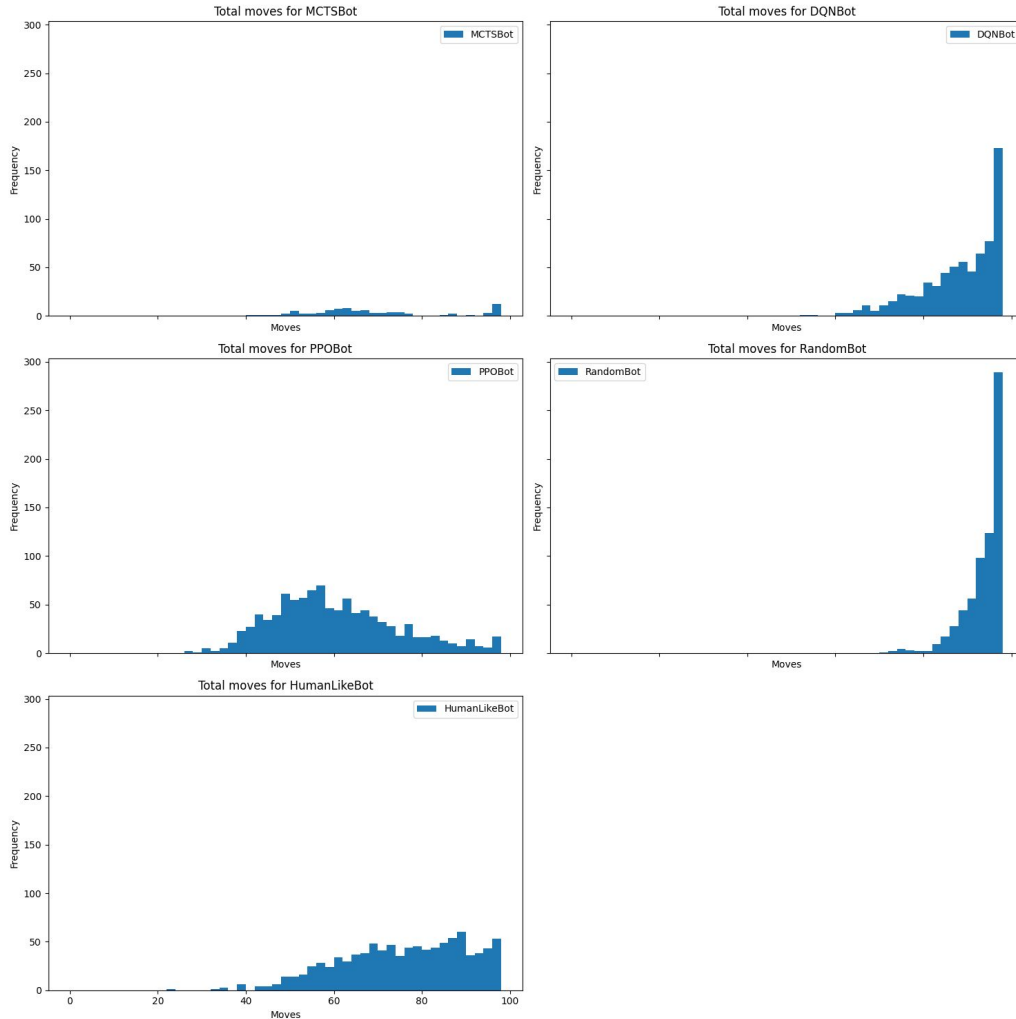


# Results

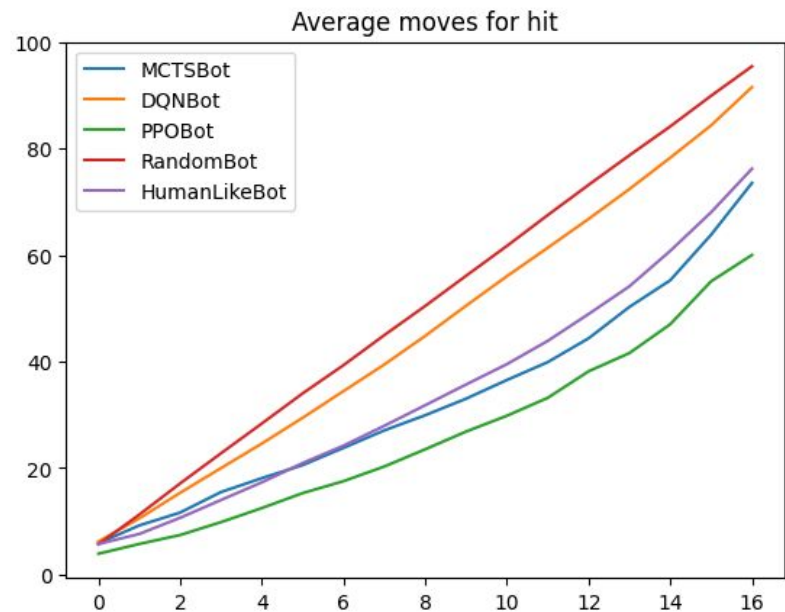

We simulated each bot for a 1000 games on random boards till it won (mcts for 100 due to time constraints):

Bot	Mean Moves	Median Moves	Max Moves	Min Moves
Random	95.47	97	100	70
Human-Like	76.23	77.5	100	22
MCTS (100)	73.55	67.5	100	41
DQN	91.55	95.0	100	52
PPO	<u>60.02</u>	<u>58.0</u>	<u>100</u>	<u>26</u>

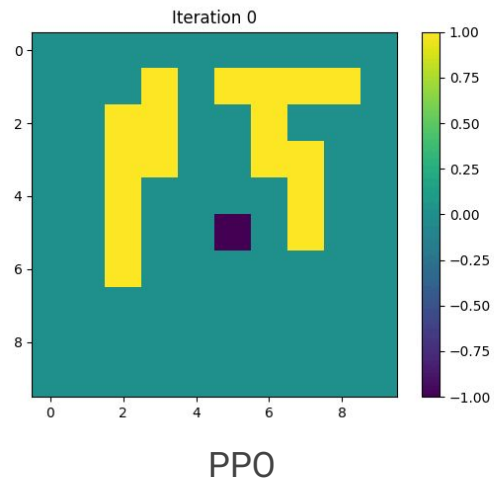
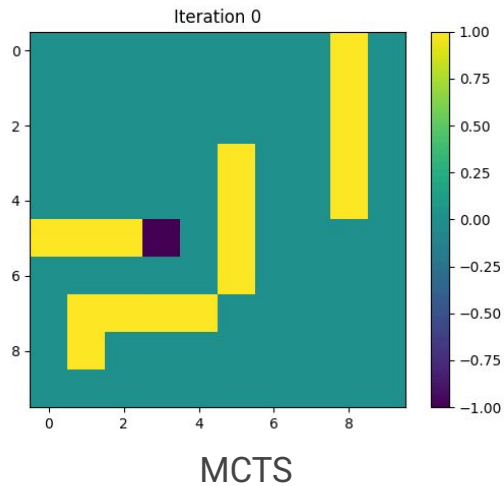
# Results



# Results



# Sample Runs



# Head to Head

We made each bot play the others for 1000 games:

	Random	Human-Like	MCTS (100)	DQN	PPO
Random	X	<u>90.4% / 74</u>	<u>85% / 69.2</u>	<u>62% / 88.14</u>	<u>98% / 59.53</u>
Human-Like	9.6% / 91	X	<u>58% / 59.00</u>	19% / 82.71	<u>80.4% / 57</u>
MCTS (100)	15% / 92.60	42% / 73.40	X	26% / 84.88	<u>57% / 56.1</u>
DQN	38% / 91.92	<u>81% / 74.06</u>	<u>74% / 65.90</u>	X	<u>96% / 59.60</u>
PPO	2% / 90	19.6% / 63	43% / 57.86	4% / 78.43	X



# Insights

- We learnt about three new algorithms
- It is difficult to train to play partially observable processes
- Battleship is a relatively small game where the winner still has a high degree of luck on their side due to randomness
- Relatively simpler methods like MCTS could perform comparably to Deep RL methods



# Demo

Play the PPO-bot!





Thank You